

# Ganeti Setup & Walk-thru Guide

- 3-node cluster using Vagrant & VirtualBox
- All setup using puppet
- Need to install git, VirtualBox, & Vagrant

# Repo & Vagrant Setup

Make sure you have hardware virtualization enabled in your BIOS prior to running VirtualBox. You will get an error from VirtualBox while starting the VM if you don't it enabled.

```
gem install vagrant  
git clone git://github.com/ramereth/vagrant-ganeti.git  
git submodule update --init
```

# Starting up & accessing the nodes

The Vagrantfile is setup to where you can deploy one, two, or three nodes depending on your use case. **Node1** will have Ganeti already initialized while the other two will only have Ganeti installed and primed. For more information on how to use Vagrant, please check out their site. *NOTE: Root password is 'vagrant' on all nodes.*

## # Starting a single node (node1)

```
vagrant up node1
```

```
vagrant ssh node1
```

## # Starting node2

```
vagrant up node2
```

```
vagrant ssh node1
```

```
gnt-node add -s 33.33.34.12 node2
```

## # Starting node3

```
vagrant up node3
```

```
vagrant ssh node1
```

```
gnt-node add -s 33.33.34.13 node3
```

# What Vagrant will do for you

1. Install all dependencies required for Ganeti
2. Setup the machine to function as a Ganeti node
3. Install Ganeti, Ganeti Htools, and Ganeti Instance Image
4. Setup and initialize Ganeti (*node1 only*)

# Installing Ganeti

We've already installed Ganeti for you on the VMs, but here are the steps that we did for documentation purposes.

```
tar -zxvf ganeti-2.5.1.tar.gz
cd ganeti-2.5.1
./configure --localstatedir=/var --sysconfdir=/etc &&
/usr/bin/make && /usr/bin/install
cp doc/examples/ganeti.initd /etc/init.d/ganeti && chmod +x
/etc/init.d/ganeti
update-rc.d ganeti defaults 20 80
```

# Initialize Ganeti

Ganeti will be already initialized on **node1** for you, but here are the steps that we did. Be aware that Ganeti is very picky about extra spaces in the “-H kvm:” line.

```
gnt-cluster init \  
  --vg-name=ganeti -s 33.33.34.11 \  
  --master-netdev=br0 \  
  -I hail \  
  -H kvm:kernel_path=/boot/vmlinuz-kvmU, \  
    initrd_path=/boot/initrd-kvmU, \  
    root_path=/dev/sda2,nic_type=e1000,disk_type=scsi, \  
    vnc_bind_address=0.0.0.0, serial_console=true \  
  -N link=br0 --enabled-hypervisors=kvm \  
ganeti.example.org
```

# Testing the cluster: verify

```
root@node1:~# gnt-cluster verify
Submitted jobs 4, 5
Waiting for job 4 ...
Thu Jun 7 06:03:54 2012 * Verifying cluster config
Thu Jun 7 06:03:54 2012 * Verifying cluster certificate files
Thu Jun 7 06:03:54 2012 * Verifying hypervisor parameters
Thu Jun 7 06:03:54 2012 * Verifying all nodes belong to an existing group
Waiting for job 5 ...
Thu Jun 7 06:03:54 2012 * Verifying group 'default'
Thu Jun 7 06:03:54 2012 * Gathering data (1 nodes)
Thu Jun 7 06:03:54 2012 * Gathering disk information (1 nodes)
Thu Jun 7 06:03:54 2012 * Verifying configuration file consistency
Thu Jun 7 06:03:54 2012 * Verifying node status
Thu Jun 7 06:03:54 2012 * Verifying instance status
Thu Jun 7 06:03:54 2012 * Verifying orphan volumes
Thu Jun 7 06:03:54 2012 * Verifying N+1 Memory redundancy
Thu Jun 7 06:03:54 2012 * Other Notes
Thu Jun 7 06:03:54 2012 * Hooks Results
```

# Testing the cluster: list nodes

```
root@node1:~# gnt-node list
```

Node	DTotal	DFree	MTotal	MNode	MFree	Pinst	Sinst
node1.example.org	26.0G	25.5G	744M	186M	587M	0	0
node2.example.org	26.0G	25.5G	744M	116M	650M	0	0



# Adding an Instance

```
root@node1:~# gnt-os list
```

```
Name
```

```
image+cirros
```

```
image+default
```

```
root@node1:~# gnt-instance add -n node1 -o image+cirros -t plain -s 1G \  
--no-start instance1
```

```
Thu Jun  7 06:05:58 2012 * disk 0, vg ganeti, name 780af428-3942-4fa9-8307-1323de416519.disk0
```

```
Thu Jun  7 06:05:58 2012 * creating instance disks...
```

```
Thu Jun  7 06:05:58 2012 adding instance instance1.example.org to cluster config
```

```
Thu Jun  7 06:05:58 2012 - INFO: Waiting for instance instance1.example.org to sync disks.
```

```
Thu Jun  7 06:05:58 2012 - INFO: Instance instance1.example.org's disks are in sync.
```

```
Thu Jun  7 06:05:58 2012 * running the instance OS create scripts...
```

# Listing Instance Information

```
root@node1:~# gnt-instance list
```

Instance	Hypervisor	OS	Primary_node	Status	Memory
instance1.example.org	kvm	image+cirros	node1.example.org	ADMIN_down	-

# Listing Instance Information

```
root@node1:~# gnt-instance info instance1
Instance name: instance1.example.org
UUID: bb87da5b-05f9-4dd6-9bc9-48592c1e091f
Serial number: 1
Creation time: 2012-06-07 06:05:58
Modification time: 2012-06-07 06:05:58
State: configured to be down, actual state is down
  Nodes:
    - primary: node1.example.org
    - secondaries:
Operating system: image+cirros
Allocated network port: 11000
Hypervisor: kvm
  - console connection: vnc to node1.example.org:11000 (display 5100)
...
Hardware:
  - VCPUs: 1
  - memory: 128MiB
  - NICs:
    - nic/0: MAC: aa:00:00:dd:ac:db, IP: None, mode: bridged, link: br0
Disk template: plain
Disks:
  - disk/0: lvm, size 1.0G
    access mode: rw
    logical_id: ganeti/780af428-3942-4fa9-8307-1323de416519.disk0
    on primary: /dev/ganeti/780af428-3942-4fa9-8307-1323de416519.disk0 (252:1)
```

# Controlling Instances

```
root@node1:~# gnt-instance start instance1  
Waiting for job 10 for instance1.example.org ...
```

```
root@node1:~# gnt-instance console instance1
```

```
login as 'vagrant' user. default password: 'vagrant'. use 'sudo' for root.  
cirros login:
```

```
# Press ctrl+] to escape console.
```

```
root@node1:~# gnt-instance shutdown instance1  
Waiting for job 11 for instance1.example.org ...
```

# Changing the Disk Type

```
root@node1:~# gnt-instance shutdown instance1
Waiting for job 11 for instance1.example.org ...
root@node1:~# gnt-instance modify -t drbd -n node2 instance1
Thu Jun  7 06:09:07 2012 Converting template to drbd
Thu Jun  7 06:09:08 2012 Creating additional volumes...
Thu Jun  7 06:09:08 2012 Renaming original volumes...
Thu Jun  7 06:09:08 2012 Initializing DRBD devices...
Thu Jun  7 06:09:09 2012 - INFO: Waiting for instance instance1.example.org to sync disks.
Thu Jun  7 06:09:11 2012 - INFO: - device disk/0:  5.10% done, 20s remaining (estimated)
Thu Jun  7 06:09:31 2012 - INFO: - device disk/0: 86.00% done,  3s remaining (estimated)
Thu Jun  7 06:09:34 2012 - INFO: - device disk/0: 98.10% done,  0s remaining (estimated)
Thu Jun  7 06:09:34 2012 - INFO: Instance instance1.example.org's disks are in sync.
Modified instance instance1
- disk_template -> drbd
Please don't forget that most parameters take effect only at the next start of the instance.
```

# Instance Failover

```
root@node1:~# gnt-instance failover -f instancel
Thu Jun  7 06:10:09 2012 - INFO: Not checking memory on the secondary node as instance will not
be started
Thu Jun  7 06:10:09 2012 Failover instance instancel.example.org
Thu Jun  7 06:10:09 2012 * not checking disk consistency as instance is not running
Thu Jun  7 06:10:09 2012 * shutting down instance on source node
Thu Jun  7 06:10:09 2012 * deactivating the instance's disks on source node
```

# Instance Migration

```
root@node1:~# gnt-instance start instancel
Waiting for job 14 for instancel.example.org ...
root@node1:~# gnt-instance migrate -f instancel
Thu Jun 7 06:10:38 2012 Migrating instance instancel.example.org
Thu Jun 7 06:10:38 2012 * checking disk consistency between source and target
Thu Jun 7 06:10:38 2012 * switching node node1.example.org to secondary mode
Thu Jun 7 06:10:38 2012 * changing into standalone mode
Thu Jun 7 06:10:38 2012 * changing disks into dual-master mode
Thu Jun 7 06:10:39 2012 * wait until resync is done
Thu Jun 7 06:10:39 2012 * preparing node1.example.org to accept the instance
Thu Jun 7 06:10:39 2012 * migrating instance to node1.example.org
Thu Jun 7 06:10:44 2012 * switching node node2.example.org to secondary mode
Thu Jun 7 06:10:44 2012 * wait until resync is done
Thu Jun 7 06:10:44 2012 * changing into standalone mode
Thu Jun 7 06:10:45 2012 * changing disks into single-master mode
Thu Jun 7 06:10:46 2012 * wait until resync is done
Thu Jun 7 06:10:46 2012 * done
```

# Master Failover

```
root@node2:~# gnt-cluster master-failover
```

```
root@node2:~# gnt-cluster getmaster  
node2.example.org
```

```
root@node1:~# gnt-cluster master-failover
```



# Job Operations

```
root@node1:~# gnt-job list
ID Status Summary
1 success CLUSTER_POST_INIT
2 success CLUSTER_SET_PARAMS
3 success CLUSTER_VERIFY
4 success CLUSTER_VERIFY_CONFIG
5 success CLUSTER_VERIFY_GROUP(8e97b380-3d86-4d3f-a1c5-c7276edb8846)
6 success NODE_ADD(node2.example.org)
7 success OS_DIAGNOSE
8 success INSTANCE_CREATE(instance1.example.org)
9 success INSTANCE_QUERY_DATA
10 success INSTANCE_STARTUP(instance1.example.org)
11 success INSTANCE_SHUTDOWN(instance1.example.org)
12 success INSTANCE_SET_PARAMS(instance1.example.org)
13 success INSTANCE_FAILOVER(instance1.example.org)
14 success INSTANCE_STARTUP(instance1.example.org)
15 success INSTANCE_MIGRATE(instance1.example.org)
```

# Job Operations

```
root@node1:~# gnt-job info 14
Job ID: 14
Status: success
Received:      2012-06-07 06:10:29.032216
Processing start: 2012-06-07 06:10:29.100896 (delta 0.068680s)
Processing end:   2012-06-07 06:10:30.759979 (delta 1.659083s)
Total processing time: 1.727763 seconds
Opcodes:
  OP_INSTANCE_STARTUP
    Status: success
    Processing start: 2012-06-07 06:10:29.100896
    Execution start:  2012-06-07 06:10:29.173253
    Processing end:    2012-06-07 06:10:30.759952
    Input fields:
      beparams: {}
      comment: None
      debug_level: 0
      depends: None
      dry_run: False
      force: False
      hvparams: {}
      ignore_offline_nodes: False
      instance_name: instance1.example.org
      no_remember: False
      priority: 0
      startup_paused: False
    No output data
    Execution log:
```

# Using Htools

```
root@node1:~# gnt-instance add -I hail -o image+cirros -t drbd -s 1G --no-start instance2
Thu Jun  7 06:14:05 2012 - INFO: Selected nodes for instance instance2.example.org via
iallocator hail: node2.example.org, node1.example.org
Thu Jun  7 06:14:06 2012 * creating instance disks...
Thu Jun  7 06:14:08 2012 adding instance instance2.example.org to cluster config
Thu Jun  7 06:14:08 2012 - INFO: Waiting for instance instance2.example.org to sync disks.
Thu Jun  7 06:14:09 2012 - INFO: - device disk/0:  6.30% done, 16s remaining (estimated)
Thu Jun  7 06:14:26 2012 - INFO: - device disk/0: 73.20% done, 6s remaining (estimated)
Thu Jun  7 06:14:33 2012 - INFO: - device disk/0: 100.00% done, 0s remaining (estimated)
Thu Jun  7 06:14:33 2012 - INFO: Instance instance2.example.org's disks are in sync.
Thu Jun  7 06:14:33 2012 * running the instance OS create scripts...
```

```
root@node1:~# gnt-instance list
```

Instance	Hypervisor	OS	Primary_node	Status	Memory
instance1.example.org	kvm	image+cirros	node1.example.org	running	128M
instance2.example.org	kvm	image+cirros	node2.example.org	ADMIN_down	-

# Using Htools: hbal

```
root@node1:~# hbal -L
Loaded 2 nodes, 2 instances
Group size 2 nodes, 2 instances
Selected node group: default
Initial check done: 0 bad nodes, 0 bad instances.
Initial score: 3.89180108
Trying to minimize the CV...
  1. instance1 node1:node2 => node2:node1 0.04771505 a=f
Cluster score improved from 3.89180108 to 0.04771505
Solution length=1
```

```
root@node1:~# hbal -L -X
Loaded 2 nodes, 2 instances
Group size 2 nodes, 2 instances
Selected node group: default
Initial check done: 0 bad nodes, 0 bad instances.
Initial score: 3.89314516
Trying to minimize the CV...
  1. instance1 node1:node2 => node2:node1 0.04905914 a=f
Cluster score improved from 3.89314516 to 0.04905914
Solution length=1
Executing jobset for instances instance1.example.org
Got job IDs 18
```

# Using Htools: hspace

```
root@node1:~# hspace --memory 128 --disk 1024 -L
The cluster has 2 nodes and the following resources:
  MEM 1488, DSK 53304, CPU 4, VCPU 256.
There are 2 initial instances on the cluster.
Normal (fixed-size) instance spec is:
  MEM 128, DSK 1024, CPU 1, using disk template 'drbd'.
Normal (fixed-size) allocation results:
- 2 instances allocated
- most likely failure reason: FailMem
- initial cluster score: 0.04233871
- final cluster score: 0.04233871
- memory usage efficiency: 34.41%
- disk usage efficiency: 18.25%
- vcpu usage efficiency: 1.56%
```

# Recovering from a Node Failure

# Setup node3

```
gnt-node add -s 33.33.34.13 node3
```

# You should see something like the following now:

```
root@node1:~# gnt-node list
```

Node	DTotal	DFree	MTotal	MNode	MFree	Pinst	Sinst
node1.example.org	26.0G	23.3G	744M	213M	585M	1	1
node2.example.org	26.0G	23.3G	744M	247M	542M	1	1
node3.example.org	26.0G	25.5G	744M	114M	650M	0	0

# Simulating a node failure

```
# Log out of node1
```

```
vagrant halt -f node2
```

```
# Log back into node1
```

```
gnt-cluster verify
```

```
gnt-node modify -O yes -f node2
```

```
gnt-cluster verify
```

```
gnt-node failover --ignore-consistency node2
```

```
gnt-node evacuate -I hail -s node2
```

```
gnt-cluster verify
```

# Re-adding node2

*# Log out of node1*

```
vagrant destroy -f node2
```

```
vagrant up node2
```

*# Log back into node1*

```
gnt-node add --readd node2
```

```
gnt-cluster verify
```